

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
```

In [2]:

```
a=pd.read_csv(r"C:\Users\Gowthami\Downloads\bottle.csv.zip")  
a
```

C:\Users\Gowthami\AppData\Local\Temp\ipykernel\_22952\1376288322.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low\_memory=False.

```
a=pd.read_csv(r"C:\Users\Gowthami\Downloads\bottle.csv.zip")
```

Out[2]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600

In [3]:

```
a.head()
```

Out[3]:

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	19	10.450	33.4200	NaN	25.64300
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0019A-3	0	10.50	33.440	NaN	25.649
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	20	10.450	33.4210	NaN	25.64300
1	...	...	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.46	33.440	NaN	25.656
864858	34404	864859	093.4 026.4	20-1611SR-MX-310-2239-09340264-0000A-7	0	18.744	33.4083	5.805	23.87055
2	1	3	054.0 056.0	20-1611SR-MX-310-2239-09340264-0000A-7	10	10.46	33.437	NaN	25.654
864859	34404	864860	093.4 026.4	20-1611SR-MX-310-2239-09340264-0002A-3	2	18.744	33.4083	5.805	23.87072
3	1	4	054.0 056.0	20-1611SR-MX-310-2239-09340264-0002A-3	19	10.45	33.420	NaN	25.643
864860	34404	864861	093.4 026.4	20-1611SR-MX-310-2239-09340264-0005A-3	5	18.692	33.4150	5.796	23.88911
4	1	5	054.0 056.0	20-1611SR-MX-310-2239-09340264-0005A-3	20	10.45	33.421	NaN	25.643
864861	34404	864862	093.4 026.4	20-1611SR-MX-310-2239-09340264-0007A-7	10	18.161	33.4062	5.816	24.01426

In [4]: Cst\_Cnt Btl\_Cnt Sta\_ID Depth\_ID Depthm T\_degC Salnty O2ml\_L STheta

a.info()

				20-					
				1611SR-					
864862	34404	864863	093.4	MX-310-	15	17.533	33.3880	5.774	24.15297
			026.4	2239-					
				09340264-					
				0015A-3					

864863 rows × 74 columns

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 864863 entries, 0 to 864862
```

```
Data columns (total 74 columns):
```

#	Column	Non-Null Count	Dtype
0	Cst_Cnt	864863 non-null	int64
1	Btl_Cnt	864863 non-null	int64
2	Sta_ID	864863 non-null	object
3	Depth_ID	864863 non-null	object
4	Depthm	864863 non-null	int64
5	T_degC	853900 non-null	float64
6	Salnty	817509 non-null	float64
7	O2ml_L	696201 non-null	float64
8	STheta	812174 non-null	float64
9	O2Sat	661274 non-null	float64
10	Oxy_μmol/Kg	661268 non-null	float64
11	BtlNum	118667 non-null	float64
12	RecInd	864863 non-null	int64
13	T_prec	853900 non-null	float64
14	T_qual	23127 non-null	float64
15	S_prec	817509 non-null	float64
16	S_qual	74914 non-null	float64
17	P_qual	673755 non-null	float64
18	O_qual	184676 non-null	float64
19	SThtaq	65823 non-null	float64
20	O2Satq	217797 non-null	float64
21	ChlorA	225272 non-null	float64
22	Chlqua	639166 non-null	float64
23	Phaeop	225271 non-null	float64
24	Phaqua	639170 non-null	float64
25	PO4uM	413317 non-null	float64
26	PO4q	451786 non-null	float64
27	SiO3uM	354091 non-null	float64
28	SiO3qu	510866 non-null	float64
29	NO2uM	337576 non-null	float64
30	NO2q	529474 non-null	float64
31	NO3uM	337403 non-null	float64
32	NO3q	529933 non-null	float64
33	NH3uM	64962 non-null	float64
34	NH3q	808299 non-null	float64
35	C14As1	14432 non-null	float64
36	C14A1p	12760 non-null	float64
37	C14A1q	848605 non-null	float64
38	C14As2	14414 non-null	float64
39	C14A2p	12742 non-null	float64
40	C14A2q	848623 non-null	float64
41	DarkAs	22649 non-null	float64
42	DarkAp	20457 non-null	float64
43	DarkAq	840440 non-null	float64
44	MeanAs	22650 non-null	float64
45	MeanAp	20457 non-null	float64
46	MeanAq	840439 non-null	float64
47	IncTim	14437 non-null	object
48	LightP	18651 non-null	float64
49	R_Depth	864863 non-null	float64
50	R_TEMP	853900 non-null	float64
51	R_POTEMP	818816 non-null	float64
52	R_SALINITY	817509 non-null	float64
53	R_SIGMA	812007 non-null	float64
54	R_SVA	812092 non-null	float64
55	R_DYNHT	818206 non-null	float64

56	R_O2	696201	non-null	float64
57	R_O2Sat	666448	non-null	float64
58	R_SIO3	354099	non-null	float64
59	R_PO4	413325	non-null	float64
60	R_NO3	337411	non-null	float64
61	R_NO2	337584	non-null	float64
62	R_NH4	64982	non-null	float64
63	R_CHLA	225276	non-null	float64
64	R_PHAEO	225275	non-null	float64
65	R_PRES	864863	non-null	int64
66	R_SAMP	122006	non-null	float64
67	DIC1	1999	non-null	float64
68	DIC2	224	non-null	float64
69	TA1	2084	non-null	float64
70	TA2	234	non-null	float64
71	pH2	10	non-null	float64
72	pH1	84	non-null	float64
73	DIC Quality Comment	55	non-null	object

dtypes: float64(65), int64(5), object(4)

memory usage: 488.3+ MB

In [5]:

```
a.tail()
```

Out[5]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta
864858	34404	864859	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0000A-7	0	18.744	33.4083	5.805	23.87055
864859	34404	864860	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0002A-3	2	18.744	33.4083	5.805	23.87072
864860	34404	864861	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0005A-3	5	18.692	33.4150	5.796	23.88911
864861	34404	864862	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0010A-3	10	18.161	33.4062	5.816	24.01426
864862	34404	864863	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0015A-3	15	17.533	33.3880	5.774	24.15297

5 rows × 74 columns



In [6]:

```
a.describe()
```

Out[6]:

	Cst_Cnt	Btl_Cnt	Depthm	T_degC	Salnty	O
count	864863.000000	864863.000000	864863.000000	853900.000000	817509.000000	696201.0
mean	17138.790958	432432.000000	226.831951	10.799677	33.840350	3.3
std	10240.949817	249664.587269	316.050259	4.243825	0.461843	2.0
min	1.000000	1.000000	0.000000	1.440000	28.431000	-0.0
25%	8269.000000	216216.500000	46.000000	7.680000	33.488000	1.3
50%	16848.000000	432432.000000	125.000000	10.060000	33.863000	3.4
75%	26557.000000	648647.500000	300.000000	13.880000	34.196900	5.5
max	34404.000000	864863.000000	5351.000000	31.140000	37.034000	11.1

8 rows × 70 columns



In [7]:

```
a.isna().any()
```

Out[7]:

```
Cst_Cnt          False
Btl_Cnt          False
Sta_ID           False
Depth_ID         False
Depthm           False
...
TA1              True
TA2              True
pH2              True
pH1              True
DIC Quality Comment  True
Length: 74, dtype: bool
```



In [8]:

```
a.isnull().sum()
```

Out[8]:

```
Cst_Cnt      0
Btl_Cnt      0
Sta_ID       0
Depth_ID     0
Depthm       0
...
TA1          862779
TA2          864629
pH2          864853
pH1          864779
DIC Quality Comment  864808
Length: 74, dtype: int64
```

In [9]:

```
a.loc[:,['Salnty','T_degC']]
```

Out[9]:

	Salnty	T_degC
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...	...	...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

In [10]:

```
a=a[['Salnty','T_degC']]
a.columns=['Sal','Temp']
```

In [11]:

```
a.head(20)
```

Out[11]:


	Sal	Temp
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45
5	33.431	10.45
6	33.440	10.45
7	33.424	10.24
8	33.420	10.06
9	33.494	9.86
10	33.510	9.83
11	33.580	9.67
12	33.640	9.50
13	33.689	9.32
14	33.847	8.76
15	33.860	8.71
16	33.876	8.53
17	NaN	8.45
18	33.926	8.26
19	33.980	7.96

In [12]:

```
sns.lmplot(x='Sal',y='Temp',data=a,order=2,ci=None)
```

Out[12]:

<seaborn.axisgrid.FacetGrid at 0x17ff59ebc10>



In [13]:

```
a.fillna(method='ffill')
```

Out[13]:

	Sal	Temp
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...	...	...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

In [14]:

```
a.fillna(value=0,inplace=True)
```

C:\Users\Gowthami\AppData\Local\Temp\ipykernel\_22952\3015722605.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
a.fillna(value=0,inplace=True)
```

In [15]:

```
x=np.array(a['Sal']).reshape(-1,1)
y=np.array(a['Temp']).reshape(-1,1)
```

In [16]:

```
a.isna().any()
```

Out[16]:

```
Sal    False
Temp   False
dtype: bool
```

In [17]:

```
a.isnull().sum()
```

Out[17]:

```
Sal      0
Temp      0
dtype: int64
```

In [18]:

```
a.dropna(inplace=True)
```

C:\Users\Gowthami\AppData\Local\Temp\ipykernel\_22952\2317726482.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
a.dropna(inplace=True)
```

In [19]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```


In [20]:

```
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```

```
0.014809436284115463
```

In [21]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_pred,color='k')  
plt.show()
```




In [22]:

```
a500=a[:, :500]  
sns.lmplot(x='Sal', y='Temp', data=a500, order=1, ci=None)
```

Out[22]:


<seaborn.axisgrid.FacetGrid at 0x17ff475fd90>



In [23]:

```
a500.fillna(method='ffill',inplace=True)
x=np.array(a500['Sal']).reshape(-1,1)
y=np.array(a500['Temp']).reshape(-1,1)
a500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.022412040359021446



In [24]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.022412040359021446

#conclusion: Linear regression is best fit for the model




## Ridge and Lasso regression

In [25]:

```
plt.figure(figsize = (10, 10))  
sns.heatmap(a.corr(), annot = True)
```

Out[25]:

<Axes: >



In [26]:

```

features = a.columns[0:2]
target = a.columns[-1]
#X and y values
X = a[features].values
y = a[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

The dimension of X\_train is (605404, 2)  
The dimension of X\_test is (259459, 2)

In [27]:

```

#Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))

```

Linear Regression Model:

The train score for lr model is 1.0  
The test score for lr model is 1.0

In [28]:

```

#Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))

```

The train score for ridge model is 0.9999999981135502  
The train score for ridge model is 0.99999999811206

In [29]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```


Ridge Model:

The train score for ridge model is 0.99999999723243

The test score for ridge model is 0.999999997231402

In [30]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



In [31]:

```
#Using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.9999999994806811

0.9999999994806712

In [32]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0


The test score for ls model is -1.9031696447013857e-05

In [33]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[33]:

<Axes: >




In [34]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',)
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
```

Out[34]:

```
[<matplotlib.lines.Line2D at 0x17fd2e42b10>]
```



## Elastic net

In [35]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[0.          0.94934511]
0.5401219631063316
```

In [36]:

```
y_pred_elastic=regr.predict(X_train)
```

In [37]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 114.40984808660129
```

## Vehicle selection

In [38]:

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge

df=pd.read_csv(r"C:\Users\Gowthami\Downloads\fiat500_VehicleSelection_Dataset.csv")
df

```

Out[38]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns





In [39]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ID                    1538 non-null   int64  
1   model                 1538 non-null   object  
2   engine_power          1538 non-null   int64  
3   age_in_days           1538 non-null   int64  
4   km                    1538 non-null   int64  
5   previous_owners       1538 non-null   int64  
6   lat                   1538 non-null   float64 
7   lon                   1538 non-null   float64 
8   price                 1538 non-null   int64  
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [40]:

```
df.head()
```

Out[40]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	15380
1	2	pop	51	1186	32500	1	45.666359	12.241890	15370
2	3	sport	74	4658	142228	1	45.503300	11.417840	15360
3	4	lounge	51	2739	160000	1	40.633171	17.634609	15350
4	5	pop	73	3074	106880	1	41.903221	12.495650	15340

In [41]:

```
df.tail()
```

Out[41]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
1533	1534	sport	51	3712	115280	1	45.069679	7.7041	15330
1534	1535	lounge	74	3835	112000	1	45.845692	8.6661	15320
1535	1536	pop	51	2223	60457	1	45.481541	9.4131	15310
1536	1537	lounge	51	2557	80750	1	45.000702	7.6821	15300
1537	1538	pop	51	1766	54276	1	40.323410	17.5681	15290

In [42]:

```
df.info
```

Out[42]:

```
<bound method DataFrame.info of
ays      km  previous_owners
0         1  lounge          51
1         2    pop          51
2         3  sport          74
3         4  lounge          51
4         5    pop          73
...      ...      ...
1533    1534  sport          51
1534    1535  lounge          74
1535    1536    pop          51
1536    1537  lounge          51
1537    1538    pop          51

      ID  model  engine_power  age_in_d
0      882   25000           1  \
1     1186   32500           1
2     4658  142228           1
3     2739  160000           1
4     3074  106880           1
...     ...     ...         ...
1533   3712  115280           1
1534   3835  112000           1
1535   2223   60457           1
1536   2557   80750           1
1537   1766   54276           1

      lat      lon  price
0  44.907242  8.611560  8900
1  45.666359  12.241890  8800
2  45.503300  11.417840  4200
3  40.633171  17.634609  6000
4  41.903221  12.495650  5700
...     ...     ...     ...
1533  45.069679  7.704920  5200
1534  45.845692  8.666870  4600
1535  45.481541  9.413480  7500
1536  45.000702  7.682270  5990
1537  40.323410  17.568270  7900

[1538 rows x 9 columns]>
```

In [43]:

```
df.describe()
```

Out[43]:

	ID	engine_power	age_in_days	km	previous_owners	li
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.54136
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.13351
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.85583
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.80295
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.39405
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.46796
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.79561

In [44]:

```
df.isna().any()
```

Out[44]:

ID	False
model	False
engine_power	False
age_in_days	False
km	False
previous_owners	False
lat	False
lon	False
price	False
dtype: bool	

In [45]:

```
df.isnull().sum()
```

Out[45]:

ID	0
model	0
engine_power	0
age_in_days	0
km	0
previous_owners	0
lat	0
lon	0
price	0
dtype: int64	

In [46]:

```
df.isnull()
```

Out[46]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...
1533	False	False	False	False	False	False	False	False	False
1534	False	False	False	False	False	False	False	False	False
1535	False	False	False	False	False	False	False	False	False
1536	False	False	False	False	False	False	False	False	False
1537	False	False	False	False	False	False	False	False	False

1538 rows × 9 columns



In [47]:

```
df.loc[:11,["ID","price"]]
```

Out[47]:

	ID	price
0	1	8900
1	2	8800
2	3	4200
3	4	6000
4	5	5700
5	6	7900
6	7	10750
7	8	9190
8	9	5600
9	10	6000
10	11	8950
11	12	10990

In [48]:


```
df=df[["engine_power","price"]]  
df.columns=["Engine","pri"]
```

In [49]:

```
sns.lmplot(x='Engine',y='pri',data=df,order=2,ci=None)
```

Out[49]:

<seaborn.axisgrid.FacetGrid at 0x17fd2e74350>



In [50]:

```
df.describe()
```

Out[50]:

	Engine	pri
count	1538.000000	1538.000000
mean	51.904421	8576.003901
std	3.988023	1939.958641
min	51.000000	2500.000000
25%	51.000000	7122.500000
50%	51.000000	9000.000000
75%	51.000000	10000.000000
max	77.000000	11100.000000

In [51]:

```
df.fillna(method="ffill")
```

Out[51]:

	Engine	pri
0	51	8900
1	51	8800
2	74	4200
3	51	6000
4	73	5700
...	...	...
1533	51	5200
1534	74	4600
1535	51	7500
1536	51	5990
1537	51	7900

1538 rows × 2 columns

In [52]:

```
x=np.array(df['Engine']).reshape(-1,1)  
y=np.array(df['pri']).reshape(-1,1)
```

In [53]:

```
df.dropna(inplace=True)
```

C:\Users\Gowthami\AppData\Local\Temp\ipykernel\_22952\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.dropna(inplace=True)
```

In [54]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```


In [55]:

```
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

0.046186912889935816

In [56]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_pred,color='k')  
plt.show()
```




In [57]:

```
df500=df[:][:500]
sns.lmplot(x='Engine',y='pri',data=df500,order=1,ci=None)
```

Out[57]:

<seaborn.axisgrid.FacetGrid at 0x17fd2e13710>



```
#df500.fillna(method='ffill',inplace=True) x=np.array(df500['Engine']).reshape(-1,1)
y=np.array(df500['pri']).reshape(-1,1) df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25) regr=LinearRegression()
regr.fit(x_train,y_train) print("Regression:",regr.score(x_test,y_test)) y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b') plt.plot(x_test,y_pred,color='k') plt.show()
```

In [58]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.046186912889935816



#conclusion: Linear regression is not fit for the model

## Ridge and Lasso regression

In [59]:

```
features = df.columns[0:2]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (1076, 2)

The dimension of X\_test is (462, 2)

In [60]:

```
#Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

In [61]:

```
#Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.9999999999999897

The train score for ridge model is 0.9999999999999899

In [62]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train, y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9999088581979684


The test score for ridge model is 0.9999100853681022

In [63]:

```
plt.figure(figsize = (10, 10))  
sns.heatmap(df.corr(), annot = True)
```


Out[63]:

<Axes: >



In [64]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



In [65]:

```
#Using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.9999999999501757

0.9999999999638806

In [66]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999728562194999


The test score for ls model is 0.9999728508562553

In [67]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[67]:

<Axes: >




In [68]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',)
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
```

Out[68]:

[<matplotlib.lines.Line2D at 0x17fd2c63290>]



## Elastic net

In [69]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[-0.          0.99999973]
0.002280249860632466
```

In [70]:

```
y_pred_elastic=regr.predict(X_train)
```

In [71]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 77371869.93693778

## Data

In [72]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
```

In [73]:

```
de=pd.read_csv(r"C:\Users\Gowthami\Downloads\data.csv")
de
```

Out[73]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0
...	...	...	...	...	...	...	...	...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0

4600 rows × 18 columns

In [74]:

```
de=de[["sqft_living","sqft_lot"]]
de.columns=["Liv","Lot"]
```




In [75]:

```
sns.lmplot(x='Liv',y='Lot',data=de,order=2,ci=None)
```

Out[75]:

<seaborn.axisgrid.FacetGrid at 0x17fd2c2f7d0>



In [76]:

```
de.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    Liv     4600 non-null   int64  
 1    Lot     4600 non-null   int64  
dtypes: int64(2)
memory usage: 72.0 KB
```

In [77]:

```
de.describe()
```

Out[77]:

	Liv	Lot
<b>count</b>	4600.000000	4.600000e+03
<b>mean</b>	2139.346957	1.485252e+04
<b>std</b>	963.206916	3.588444e+04
<b>min</b>	370.000000	6.380000e+02
<b>25%</b>	1460.000000	5.000750e+03
<b>50%</b>	1980.000000	7.683000e+03
<b>75%</b>	2620.000000	1.100125e+04
<b>max</b>	13540.000000	1.074218e+06

In [78]:

```
de.fillna(method='ffill')
```

Out[78]:

	Liv	Lot
<b>0</b>	1340	7912
<b>1</b>	3650	9050
<b>2</b>	1930	11947
<b>3</b>	2000	8030
<b>4</b>	1940	10500
...	...	...
<b>4595</b>	1510	6360
<b>4596</b>	1460	7573
<b>4597</b>	3010	7014
<b>4598</b>	2090	6630
<b>4599</b>	1490	8102

4600 rows × 2 columns

In [79]:

```
x=np.array(de['Liv']).reshape(-1,1)  
y=np.array(de['Lot']).reshape(-1,1)
```

In [80]:

```
de.dropna(inplace=True)
```

C:\Users\Gowthami\AppData\Local\Temp\ipykernel\_22952\836337131.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
de.dropna(inplace=True)
```

In [81]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```


In [82]:

```
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

0.10640487967838053

In [83]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_pred,color='k')  
plt.show()
```




In [84]:

```
de500=de[:][:500]  
sns.lmplot(x='Liv',y='Lot',data=de500,order=1,ci=None)
```

Out[84]:


<seaborn.axisgrid.FacetGrid at 0x17fd2e50650>



In [85]:

```
de500.fillna(method='ffill',inplace=True)
x=np.array(de500['Liv']).reshape(-1,1)
y=np.array(de500['Lot']).reshape(-1,1)
de500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.11354254617812931



In [86]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.11354254617812931

#conclusion: Linear regression is fit for the model


## Ridge and lasso

In [87]:

```
plt.figure(figsize = (10, 10))  
sns.heatmap(de.corr(), annot = True)
```

Out[87]:

<Axes: >



In [88]:

```

features = de.columns[0:2]
target = de.columns[-1]
#X and y values
X = de[features].values
y = de[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

The dimension of X\_train is (3220, 2)

The dimension of X\_test is (1380, 2)

In [89]:

```

#Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))

```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

In [90]:

```

#Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))

```

The train score for ridge model is 0.9999999999999997

The train score for ridge model is 0.9999999999999996

In [91]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9999900245012017


The test score for ridge model is 0.9999902306741419



In [92]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')

plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



In [93]:

```
#Using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.9999999999997705

0.9999999999996928

In [94]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999999291360324


The test score for ls model is 0.9999999291231869

In [95]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[95]:

<Axes: >




In [96]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',)
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
```

Out[96]:

[<matplotlib.lines.Line2D at 0x17ffa835350>]



## Elastic net

In [97]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[7.84570096e-07  9.99999995e-01]
-0.001601077769009862
```

In [98]:

```
y_pred_elastic=regr.predict(X_train)
```

In [99]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 1635484254.240867
```

In [ ]: